

# RFC-001: Replication Layer for Library of Cyberspace

**Status:** Implemented **Date:** January 2026 **Author:** Derrell Piper ddp@eludom.net **Implementation:** vault.scm (lines 507-713)

---

## Abstract

This RFC specifies a replication layer for the Library of Cyberspace preservation architecture, enabling cryptographically sealed releases to be published, subscribed to, and synchronized across distributed locations while maintaining tamper-evident audit trails.

## Motivation

The Library of Cyberspace requires a distribution mechanism that:

1. **Preserves cryptographic authenticity** - Signatures travel with artifacts
2. **Enables offline verification** - No centralized authority required
3. **Records provenance** - All publication events are audited
4. **Supports multiple transports** - Git, HTTP, filesystem
5. **Maintains loose coupling** - Works for confederations of friends

Traditional package managers and distribution systems assume centralized registries and online verification. This replication layer is designed for decentralized, long-term preservation where trust is established through SPKI certificates and cryptographic seals.

## Design Principles

1. **Sealed Releases** - Only cryptographically signed releases can be published
2. **Transport Agnostic** - Same API works for git, HTTP, filesystem
3. **Audit Everything** - All replication events are recorded in tamper-evident log
4. **Verify Before Trust** - Subscribers must verify cryptographic seals
5. **Explicit Authorization** - SPKI certificates determine who can publish

## Specification

### Three Core Operations

1. **seal-publish** Publish a sealed release to a remote location.

```
(seal-publish version
    remote: target
    archive-format: format
    message: notes)
```

**Parameters:** - `version` - Semantic version string (e.g., “1.0.0”) - `remote` - Publication target (git remote, URL, or directory path) - `archive-format` - ‘tarball’, ‘bundle’, or ‘cryptographic’ (default) - `message` - Release notes (optional)

**Behavior:** 1. Verify release exists (creates if needed via seal-release) 2. Create cryptographic archive with: - Tarball of repository at version tag - SHA-512 hash of tarball - Ed25519 signature of hash - Manifest with version, hash, signature 3. Publish to remote based on type: - **Git remote**: Push tag, optionally upload archive - **HTTP URL**: POST archive to endpoint - **Filesystem**: Copy archive to directory 4. Record publication in audit trail with: - Actor (public key from signing key) - Action (seal-publish version remote) - Motivation (release notes) - Cryptographic seal (signature)

#### Audit Entry Format:

```
(audit-entry
  (id "sha512:...")
  (timestamp "Mon Jan 5 23:38:20 2026")
  (sequence 1)
  (actor
    (principal #${public-key-blob})
    (authorization-chain))
  (action
    (verb seal-publish)
    (object "1.0.0")
    (parameters "/path/to/remote"))
  (context
    (motivation "Published to filesystem")
    (language "en"))
  (environment
    (platform "unknown")
    (timestamp 1767685100))
  (seal
    (algorithm "ed25519-sha512")
    (content-hash "..."))
  (signature "...")))
```

**2. seal-subscribe** Subscribe to sealed releases from a remote source.

```
(seal-subscribe remote
    target: local-path
    verify-key: public-key)
```

**Parameters:** - `remote` - Source location (git remote, URL, or directory) - `target` - Local path for downloaded archives (optional) - `verify-key` - Public key for signature verification (optional)

**Behavior:** 1. Discover available releases from remote: - **Git remote:** List tags - **HTTP URL:** GET /releases endpoint - **Filesystem:** List .archive files  
2. Download cryptographic archives  
3. Verify each archive: - Check manifest structure - Verify SHA-512 hash of tarball - Verify Ed25519 signature (if verify-key provided)  
4. Extract verified archives to target directory  
5. Record subscription in audit trail: - Count of releases downloaded - Source location - Verification status

**Security Consideration:** Without verify-key, subscription downloads archives but cannot verify authenticity. SPKI certificate chains should be used to establish trust.

### 3. `seal-synchronize` Bidirectional synchronization of sealed releases.

```
(seal-synchronize remote
    direction: 'both
    verify-key: public-key)
```

**Parameters:** - `remote` - Sync target (git remote, URL, or directory) - `direction` - 'both (default), 'push-only, or 'pull-only - `verify-key` - Public key for signature verification (optional)

**Behavior:** 1. Discover local and remote releases  
2. Compare versions to determine: - Releases to push (local but not remote) - Releases to pull (remote but not local)  
3. Execute publication for new local releases  
4. Execute subscription for new remote releases  
5. Record synchronization in audit trail: - Count pushed and pulled - Remote location - Direction

**Use Case:** Periodic sync between trusted peers in a confederation.

## Archive Format

### Cryptographic Archive Structure

```
vault-1.0.0.archive          # Manifest file
vault-1.0.0.archive.tar.gz   # Actual tarball
```

### Manifest S-expression:

```
(sealed-archive
  (version "1.0.0")
  (format cryptographic)
  (tarball "vault-1.0.0.archive.tar.gz")
  (hash "sha512:...")
  (signature "ed25519:..."))
```

```
(timestamp 1767685100)
(sealer #${public-key-blob}))
```

**Verification Steps:** 1. Read manifest 2. Hash tarball with SHA-512 3. Verify hash matches manifest 4. Verify Ed25519 signature on hash 5. Check SPKI authorization (optional)

## Transport Implementations

### Git Remote

- Uses `git push` to share tags
- Optionally uploads archives as release assets (GitHub, GitLab)
- Fetch uses `git fetch + git tag -l`

### HTTP Endpoint

- POST to `/releases/` for publication
- GET `/releases` for discovery
- Content-Type: `application/x-sealed-archive`

### Filesystem

- Copy archives to shared directory
- Directory structure: `<remote>/<archive-name>`
- No network required, works with NFS, USB drives, etc.

## Audit Integration

Every replication operation creates an audit entry with:

1. **Content-addressed ID** - SHA-512 hash of entry
2. **Chained structure** - References parent entry
3. **SPKI principal** - Public key of actor
4. **Dual context** - Human motivation + machine environment
5. **Cryptographic seal** - Ed25519 signature

This provides: - **Non-repudiation** - Cannot deny publication - **Tamper evidence** - Changes are detectable - **Causality** - Chain shows temporal order - **Accountability** - Know who published what when

## Security Considerations

### Threat Model

**Trusted:** - Local filesystem and vault - SPKI private keys - Cryptographic primitives (libsodium)

**Untrusted:** - Remote repositories - Network transport - Downloaded archives - Remote publishers (until SPKI verified)

## Attack Scenarios

1. **Malicious Archive Substitution**
  - Attacker replaces archive on remote
  - **Mitigation:** Signature verification fails
2. **Version Rollback Attack**
  - Attacker removes newer releases
  - **Mitigation:** Audit trail shows previous versions
3. **Unauthorized Publication**
  - Attacker publishes fake release
  - **Mitigation:** SPKI authorization chain required
4. **Transport Tampering**
  - Network attacker modifies download
  - **Mitigation:** Hash and signature verification

## Best Practices

1. **Always verify signatures** - Use verify-key parameter
2. **Check SPKI certificates** - Verify authorization chain
3. **Maintain audit trail** - Detect suspicious patterns
4. **Use HTTPS for HTTP transport** - Prevent network attacks
5. **Backup signing keys** - Use Shamir secret sharing

## Implementation Notes

### Helper Functions

```
(tag-exists? tag-name)      ; Check if git tag exists
/git-remote? str)          ; Detect git remote format
/http-url? str)            ; Detect HTTP/HTTPS URL
(publish-filesystem remote version archive) ; Copy to directory
(publish-http url version archive)        ; POST to endpoint
```

### Dependencies

- **Git** - For version control and tag management
- **libsodium** - Ed25519 signatures, SHA-512 hashing
- **Chicken Scheme modules:**
  - `(chicken process)` - Run git commands
  - `(chicken file)` - Filesystem operations
  - `(chicken iregex)` - URL/remote detection

## Compatibility

This specification is compatible with:

- **Git tags** - Standard git operations
- **Git bundles** - Portable repository format

- **Tarball archives** - Universal archive format
- **S-expressions** - LISP/Scheme readable format
- **SPKI/SDSI** - Authorization certificates

Future extensions may add:

- **IPFS transport** - Content-addressed distribution
- **Tor hidden services** - Anonymous publication
- **Encrypted archives** - Confidential distribution
- **Multi-signature releases** - Threshold authorization

## Test Coverage

See `test-replication.scm`:

```
;; Test seal-publish to filesystem
(seal-publish "1.0.0"
  remote: "/tmp/cyberspace-publish-test"
  message: "Published to filesystem")

;; Verify archive exists
(file-exists? "/tmp/cyberspace-publish-test/vault-1.0.0.archive")

;; Verify audit entry created
(audit-read sequence: 1)
```

## References

1. **SPKI/SDSI** - RFC 2693, RFC 2692
2. **Content-Addressed Storage** - Git internals, IPFS
3. **Semantic Versioning** - semver.org
4. **Ed25519** - Bernstein et al.
5. **Audit Trails** - RFC-002 (Cryptographic Audit Trail)

## Changelog

- **2026-01-05** - Initial implementation and specification
  - seal-publish with git/HTTP/filesystem support
  - seal-subscribe with signature verification
  - seal-synchronize with bidirectional sync
  - Full audit trail integration
  - Cryptographic archive format

---

**Implementation Status:** Complete **Test Status:** Passing **Audit Integration:** Complete